Basic Computing 2 — Packages, functions and better code[*]

Stefano Allesina and Peter Carbonetto   *University of Chicago*

The aims of this workshop are to: (1) learn how to install, load and use the many of the freely available `R` packages; (2) illustrate how to write user-defined functions, and how to organize and improve your code; use basic plotting functions; and (3) introduce the package `knitr` for writing beautiful reports. *This workshop is intended for biologists with basic knowledge of `R`.*

Setup

To follow the examples below, you will first need to install the **knitr** and **readr** packages. You will also use the **MASS** package for statistics (this package is already included with `R`).

Introductory activity: Exploring simple steps for creating better code

Much of Basic Computing 2—as well as some of the other qBio tutorials—is about the *practice* of coding in R. Using R effectively for your research projects builds on Good Practice; without it, your analyses will quickly become unmanageable, you will become discouraged, and your progress will be slow.

It may seem premature to talk about practice when we have only just started learning about R (and particularly so if you are new to programming!). But you will find that it isn't long into a project—sometimes with as little as 40–60 lines of code—when you need strategies to help you break out of a jam.

Many researchers, of course, manage to push through without adopting Good Practices. This is because they are creative and are able to perservere. But perserverance and creativity are finite resources that we would rather you use elsewhere!

Incorporating Good Practices into your work is not something that will happen overnight; it will take time and experience to develop your practice. Our aim here, in this short activity, and in other parts of Basic Computing 2, is to give you a head start—in particular, to show you that *simple steps* can go a long way to improving your coding practice.

We begin this short activity with a *script* implementing the initial steps in an exploratory analysis of tornado data from the National Weather Service. This script is `tornadoes.R`. For convenience, I've added it here:

```
# This short script produces a plot showing the number of tornado
# events recorded by the NOAA's National Weather Service on each day
# of 2011.
```

```
# Load the storm event data downloaded from
# www.ncdc.noaa.gov/stormevents/ftp.jsp
library(readr)
storms <- read_csv("storms2011.csv.gz", guess_max = 5000)
class(storms) <- "data.frame"

# Add columns for date and "day of year" (number from 1 to 365).
storms <- transform(storms, date = as.Date(paste(2011, BEGIN_YEARMONTH - 201100,
                                           BEGIN_DAY, sep = "-")))
storms <- transform(storms, dayofyear = as.numeric(format(date, "%j")))

# Focus on storm events classified as tornadoes.
tornadoes <- subset(storms, EVENT_TYPE == "Tornado")

# Plot the number of tornado events per day.
breaks      <- seq(0, 365, 3)
first_days  <- c(1,32,60,91,121,152,182,213,244,274,305,335)
months      <- c("Jan","Feb","Mar","Apr","May","Jun","Jul","Aug",
                 "Sep","Oct","Nov","Dec")
counts      <- as.numeric(table(cut(tornadoes$dayofyear, breaks)))
plot(breaks[-1], counts, type = "l", xaxt = "n")
axis(1, cex.axis = 0.7, at = first_days, labels = months)
```

If you run this code—say, by copying and pasting the code into the Console, or by running `source("tornadoes.R")`—you should see a plot showing the number of tornado events that were recorded on each day of 2011. You should also see a large "spike" in tornadoes at the end of April.

This code is quite a bit more complicated than any of the code we've seen before in Basic Computing 1. We will not spend much time trying to understand it. Before moving on, however, this is a good opportunity to point out some of the Good Practices that are implemented in this script:

1. It is self-contained—that is, it is a complete script that includes steps such as loading the necessary packages.
2. The steps are written in a logical order, starting with loading the data, and ending with a result.
3. The variable names are helpful.
4. Short, high-level comments explain what the code does.

My hope is that all these elements become part of your coding practice over time.

This script only implements the very first step in the analysis. There is still much we would like to learn, including the tornadoes' severity, and their geographic distribution. For example, are these patterns repeated in Illinois? Using `subset`, and copying some of the code above, we can investigate this question:

```
temp    <- subset(tornadoes,STATE == "ILLINOIS")
counts <- as.numeric(table(cut(temp$dayofyear, breaks)))
plot(breaks[-1], counts, type = "l", xaxt = "n")
axis(1, cex.axis = 0.7, at = first_days, labels = months)
```

This exploration will quickly become tedious if we do it in this way: fiddling with the code, and re-running it to produce a new result. You will also find that it makes difficult to keep track of the code that produced the most interesting insights.

Fortunately, there is a better way: we can *package* this code into a function. Let's begin with the part of the code that does the calculation of the daily counts (to be precise, it counts the number of events in 3-day intervals):

```
count_events <- function (dat) {
  x <- as.numeric(table(cut(dat$dayofyear, breaks)))
  return(x)
}
```

On its own, this code doesn't do anything except create the new function, `count_events`. You need to invoke it to do something useful:

```
dat    <- subset(tornadoes, STATE == "ILLINOIS")
counts <- count_events(dat)
print(counts)
```

We can now use this function anywhere we choose. For example, suppose we wanted to see whether these same patterns hold up in the more severe tornadoes (as measured by the "Fujita scale"). The `count_events` function allows us to implement this multi-step calculation in a one-liner:

```
counts <- count_events(subset(tornadoes, TOR_F_SCALE == "EF2"))
print(counts)
```

Before turning to the plots, let's take a moment to reflect on the immediate benefits of having created the `count_events` function:

1. It simplifies any parts of our analysis that make use of the this calculation.
2. It reduces the potential to make mistakes because we have wrapped these tricky calculations inside a function with a memorable name;
3. Avoids repetition of complicated code.
4. Packaging this code in a function suggests its *generality* (nothing in this code is specific to tornado events), and thereby promotes code reuse.

To further streamline our analysis, let's wrap the complicated plotting code inside a function:

```
plot_events <- function (dat) {
  x <- count_events(dat)
  plot(breaks[-1], x, type = "l", xaxt = "n")
  axis(1, cex.axis = 0.7, at = first_days, labels = months)
}
```

Let's try out this new function:

```
plot_events(tornadoes)
plot_events(subset(tornadoes, TOR_F_SCALE == "EF2"))
plot_events(subset(tornadoes, STATE == "ALABAMA"))
```

Note that this function does not output anything; the important action is to draw something to the screen.

Observe that we are making good use of our `count_events` function. Splitting the analysis into two functions seems logical; we have separated the calculation of the counts from plotting of the counts.

Again, this function is general—there is nothing about this code that is specific to tornadoes. So we may be able to reuse this function for similar analyses (e.g., studying annual flooding patterns).

Incorporating the above improvements (and a couple other small improvements), our new analysis script looks like this (see also `tornadoes_better.R`):

```r
# This short script produces a plot showing the number of tornado
# events recorded by the NOAA's National Weather Service on each day
# of 2011.

# Load the storm event data downloaded from
# www.ncdc.noaa.gov/stormevents/ftp.jsp
library(readr)
storms <- read_csv("storms2011.csv.gz", guess_max = 5000)
class(storms) <- "data.frame"

# Add columns for date and "day of year" (number from 1 to 365).
storms <- transform(storms, date = as.Date(paste(2011, BEGIN_YEARMONTH - 201100,
                                           BEGIN_DAY, sep = "-")))
storms <- transform(storms, dayofyear = as.numeric(format(date, "%j")))

# Focus on storm events classified as tornadoes.
tornadoes <- subset(storms, EVENT_TYPE == "Tornado")

# Returns a vector of length 365 giving the number of storm events per day.
count_events <- function (dat) {
  x <- as.numeric(table(cut(dat$dayofyear, breaks)))
  return(x)
}

# Plots the number of events per day.
plot_events <- function (dat, title) {
  x <- count_events(dat)
  plot(breaks[-1], x, type = "l", xaxt = "n", main = title)
  axis(1, cex.axis = 0.7, at = first_days, labels = months)
}
```

```
# Plot the total number of tornado events per day.
breaks     <- seq(0, 365, 3)
first_days <- c(1,32,60,91,121,152,182,213,244,274,305,335)
months     <- c("Jan","Feb","Mar","Apr","May","Jun","Jul","Aug",
                "Sep","Oct","Nov","Dec")
layout(matrix(1:2))
plot_events(tornadoes, "Tornadoes in 2011")

# Plot more severe tornadoes only.
plot_events(subset(tornadoes, TOR_F_SCALE == "EF2"), "Severe tornadoes in 2011")
```

As we continue to develop our analysis, can also functions reduce the effort of making revisions; for example, if we want to change the colour of the lines in all the plots, this requires only one small change to the code inside `plot_events`.

Designing functions to improve your code is very much an "art" that you will learn as you gain experience. But don't get hung up on finding the best way to organize your code into functions— it starts with recognizing the most obvious opportunities, and over time you will refine your practice.

Group activity: Compare 2011 and 2019 tornado patterns

Create a copy of `tornadoes_better.R` and call it `tornadoes_generic.R`. Modify the code in `tornadoes_generic.R` so that it can be used to analyze either the 2011 or 2019 data. The first few lines of your script should look like this:

```
year <- 2011
library(readr)
storms <- read_csv(paste0("storms",year,".csv.gz"),guess_max = 5000)
```

Hint: `paste` may be useful.

Run `tornadoes_generic.R` on the 2011 data, then re-run on the 2019 data by changing the first line to `year <- 2019`. What is the most striking difference in the 2011 and 2019 tornado patterns?

A brief tour of packages, functions, loops, and other topics

Packages

R is the most popular statistical computing software among biologists. One reason for its popularity is the availability of many packages for tackling specialized research problems. These packages are often written by biologists for biologists. You can contribute a package, too! The RStudio website (`goo.gl/harVqF`) provides guidance on how to start developing R packages. See also Hadley Wickham's free online book (`r-pkgs.had.co.nz`).

You can often find highly specialized packages to address your research questions. Here are some suggestions for finding an appropriate package. The Comprehensive R Archive Network (CRAN)

5

offers several ways to find specific packages for your task. You can browse the full list of CRAN packages (goo.gl/7oVyKC). Or you can go to the CRAN Task Views (goo.gl/0WdIcu) and browse a compilation of packages related to a topic or discipline.

From within R or RStudio, you can also call the function `RSiteSearch("keyword")`, which submits a search query to the website search.r-project.org. The website rseek.org casts an even wider net, as it not only includes package names and their documentation, but also blogs and mailing lists related to R. If your research interests are to high-throughput genomic data or other topics in bioinformatics or computational biology, you should also search the packages provided by Bioconductor (goo.gl/7dwQlq).

Installing a package

Suppose you want to install the rsvd package. The **rsvd** package provides functions to quickly perform singular value decompositions (SVD) and principal components analysis (PCA) on large data sets. To install the package, run:

```
install.packages("rsvd")
```

Or, in RStudio, select the **Packages** panel, and click **Install**.

Loading packages

Once it is successfully installed, to load the **rsvd** package into your R environment, run:

```
library(rsvd)
```

Once you have loaded the package, you can use, for example, the `rpca` function for running PCA. To access the documentation that explains what `rpca` does, and how to use it, type

```
help(rpca)
```

Now suppose you would like to access the "bacteria" data set, which reports the incidence of *H. influenzae* in Australian children. The data set is included with **MASS** package. If you try to access these data before loading the package, you will get an error:

```
data(bacteria)
```

First, you need to load the package:

```
library(MASS)
```

Now the data set is available, and you can load it:

```
data(bacteria)
head(bacteria)
```

Random numbers

Your will sometimes need to generate random numbers. (They are actually "pseudorandom" numbers because they are not perfectly random.) In fact, random numbers are needed in the "case study" below.

R has many functions to sample random numbers from different statistical distributions. For example, use `runif` to create a vector containing 10 random numbers:

```
runif(10)
```

Question: What kind of random numbers are generated by `runif`? How could you check this?

To sample from a set of values, use `sample`:

```
v <- c("a", "b", "c", "d")
sample(v, 2)                    # Sample without replacement.
sample(v, 6, replace = TRUE)    # Sample with replacement.
sample(v)                       # Shuffle the elements.
```

The normal distribution is one of the most commonly used distributions, so naturally there is a function in R for simulating from the normal distribution:

```
rnorm(3)                     # Three draws from the standard normal.
rnorm(3, mean = 5, sd = 4)   # Change the mean and standard deviation.
```

Exercise: The normal distribution has a familiar shape. Use `rnorm` to generate a large number of values from the standard normal, then use `hist` to draw a histogram of these values (you can adjust the number of bins in the histogram with the `n` argument). Is the histogram "bell shaped"? Use `mean`, `median` and `sd` to verify that the random numbers recover the expected properties of the normal distribution. Here is some code to get you started:

```
set.seed(1)
x <- rnorm(10000)
hist(x, n = 64)
```

Why is `set.seed` useful? What happens if we remove the call to `set.seed`?

Writing functions

It is good practice to subdivide your analysis into functions, and then write a short "master" program that calls the functions and performs the analysis. In this way, the code will be more legible, easier to debug, and you will be able to recycle the functions for your other projects.

In R, every function has this form:

```
my_function_name <- function (arg1, arg2, arg3) {
  #
  # Body of the function.
  #
  return(return_value) # Not required, but most functions output something.
}
```

Here is a very simple example:

```
sum_two_numbers <- function (a, b) {
  s <- a + b
  return(s)
}
sum_two_numbers(5, 7.2)
```

In R, a function can return only one object. If you need to return multiple values, organize them into a vector, matrix or list, and return that; e.g.,

```
sum_and_prod <- function (a, b) {
  s <- a + b
  p <- a * b
  return(c(s,p))
}
sum_and_prod(5, 7.2)
```

Here is a more interesting function. It accepts two arguments, x and s, and returns the density of the normal distribution with zero mean and standard deviation s at x. This is the mathematical formula for the normal density with mean zero and standard deviation s:

$$\frac{e^{-(x/s)^2/2}}{\sigma\sqrt{2\pi}}$$

Let's call this function normpdf:

```
normpdf <- function (x, s) {
  y <- exp(-(x/s)^2/2)/(sqrt(2*pi)*s)
  return(y)
}
```

Exercise: Check that this function gives the correct answers by comparing to the built-in function dnorm.

When developing and testing your function, remember this rule: *Whatever happens in the function stays in the function.* Inside normpdf, a new object, y, is created. But you will not see y in your environment (unless you happen to already have an object named y).

Vectorization

R has a feature called *vectorization*—many operations in R, including most of the basic mathematical operations, are automatically applied to all values in a vector. (We briefly learned about vectorization in Basic Computing 1.) Let's check whether R automatically vectorizes the normpdf function by running this code:

```
n <- 1000
x <- seq(-3, 3, length.out = n)
y <- normpdf(x, 1)
print(y)
plot(x, y, type = "l")
```

Activity (optional): Which operations in normpdf were applied 1,000 times, and which were applied just once? It may not be immediately obvious just by looking at the code, so to investigate this question, try running parts of the code in the console, e.g., exp(-(x/s)^2/2), and see what happens.

Vectorization is very powerful, but it may take time to get comfortable with it, and know when it will work.

Conditional branching

When we want a block of code to be executed only when a certain condition is met, we can write a conditional branching point. The syntax is as follows:

```
if (condition is met) {
  # Execute this block of code.
} else {
  # Execute this other block of code.
}
```

For example, try running these lines of code (you might want to try running them a few times):

```
x <- rnorm(1)
if (x < 0) {
  msg <- paste(x, "is less than zero")
} else if (x > 0) {
  msg <- paste(x, "is greater than zero")
} else {
  msg <- paste(x, "is equal to zero")
}
print(msg)
```

We have created a conditional branching point, so that the value of msg changes depending on whether x is less than zero, greater than zero, or equal to zero.

Activity: Improved normal probability density function

The probability density function of the normal distribution is not defined if the standard devation is less than zero. When the standard deviation is exactly zero, the density is a "spike" at zero; it is `Inf` exactly at $x = 0$, and zero everywhere else. The pseudocode for this improved normal probability density function might look something like this:

```
normpdf(x, s)
  if s < 0
    return NaN
  else if s = 0
    if x = 0
      return Inf
    else
      return 0
  else
    evaluate normal pdf with s.d. s at x
```
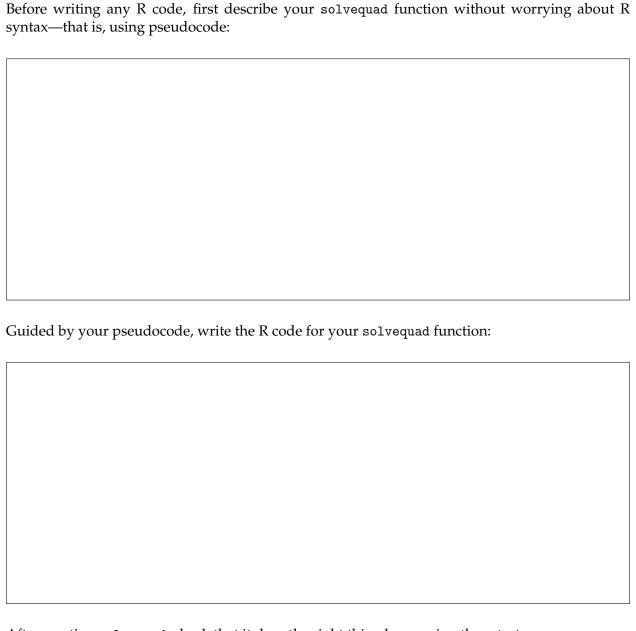
Using this pseudocode as a guide, write an improved `normpdf` function:

Activity: The quadratic formula

There is a famous formula used to solve for x in the quadratic equation $ax^2 + bx + c = 0$. It is called the *quadratic formula*. Write down the quadratic formula here:

Write a function, `solvequad`, that takes three numbers as input (a, b and c) and returns the solution(s) x that are real (*i.e.*, not complex). **Hint:** Recall that a quadratic equation may have more than one real solution—or it may have none! You will need to use `if` and `else` to handle all possible cases.

Before writing any R code, first describe your `solvequad` function without worrying about R syntax—that is, using pseudocode:

Guided by your pseudocode, write the R code for your `solvequad` function:

After creating `solvequad`, check that it does the right thing by running these tests:

```
solvequad(4, 4, 1)    # Should return -1/2.
solvequad(1, -1, -2)  # Should return 2 and -1.
solvequad(1, 1, 1)    # Should return no solutions.
```

Run a few more tests using www.wolframalpha.com.

Looping

Another way to change the flow of your program is to write a loop. A loop is simply a series of commands that are repeated a number of times. For example, you want to run the same analysis

11

on different data sets that you collected; you want to plot the results contained in a set of files; you want to test your simulation over a number of parameter sets; etc.

R provides you with two ways to loop over code blocks: the `for` loop and the `while` loop. Let's start with the `for` loop, which is used to iterate over a vector or list; for each value of the vector (or list), a series of commands will be run, as shown by the following example:

```r
v <- 1:10
for (i in v) {
  a <- i ^ 2
  print(a)
}
```

In the code above, the variable `i` takes the value of each element of `v` in sequence. Inside the block within the `for` loop, you can use the variable `i` to perform operations.

The anatomy of the `for` statement:

```r
for (variable in list_or_vector) {
  execute these commands
} # Automatically moves to the next value.
```

You should use a for loop when you know that you want to perform the analysis over a given set of values (e.g., files of a directory, rows in your data frames, sequences of a fasta file, *etc*).

The `while` loop is used when the commands need to be repeated while a certain condition is true, as shown by the following example:

```r
i <- 1
while (i <= 10) {
  a <- i ^ 2
  i <- i + 1
  print(a)
}
```

The script gives exactly the same result as the `for` loop above. A key difference is that you need to include a step to update the value of i, (using i <- i + 1), whereas in the `for` loop it is done for you automatically. The anatomy of the `while` statement:

```r
while (condition is met) {
  execute these commands
} # Beware of infinite loops... remember to update the condition!
```

You can break a loop using `break`. For example:

```r
i <- 1
while (TRUE) {
  if (i > 10) {
```

```
        break
    }
    a <- i ^ 2
    i <- i + 1
    print(a)
}
```

Question: Above, we ran three different loops, we found that each of them accomplished the same thing. Is one approach better? Why?

Programming challenge

Instructions

You will work with your group to solve the exercises below. When you have found the solutions, go to https://stefanoallesina.github.io/BSD-QBio4 and follow the link "Submit solution to challenge 2" to submit your answer (alternatively, you can go directly to goo.gl/forms/QJhKmdGqRCIuGNPa2). At the end of the bootcamp, the group with the highest number of correct solutions will be declared the winner. If you need extra time, you can work with your group during free time or in the breaks in the schedule.

Google Flu Trends

Google Flu started strong, with a paper in *Nature* (Ginsberg *et al*, 2009, doi:10.1038/nature07634) showing that, using data on Web search engine queries, one could predict the number of physician visits for influenza-like symptoms. Over time, the quality of predictions degraded considerably, requiring many adjustments to the model. Now defunct, Google Flu Trends has been proposed as a poster child of "Big Data hubris" (Lanzer *et al*, *Science*, 2014, doi:10.1126/science.1248506). In the folder containing the Basic Computing 2 tutorial materials, you will find the data used by Preis and Moat in their 2014 paper (doi:10.1098/rsos.140095) to show that, after accounting for some additional historical data, Google Flu Trends are correlated with outpatient visits due to influenza-like illnesses.

1. Read the data using function read.csv, and plot the number of weekly outpatient visits versus the Google Flu Trends estimates.
2. Calculate the (Pearson's) correlation using the cor function.
3. The data span 2010–2013. In August 2013, Google Flu changed their algorithm. Did this lead to improvements? Compare the data from August and September 2013 with the same months in 2010, 2011 and 2012. For each, calculate the correlation, and see whether the correlation is higher for 2013.

Hint: You will need to extract the year from a string for each row. This can be done using substr(gf$WeekCommencing, 1, 4), in which gf is the data frame containing the Google Flu data.

Case study: Do shorter titles lead to more citations? (optional)

To keep learning about `R`, we study the following question:

Is the length of a paper title related to the number of citations?

This is what Letchford *et al* claimed (`doi:10.1098/rsos.150266`). In 2015, they analyzed 140,000 papers, and they found that *shorter titles* were associated with a larger number of citations.

In the folder containing the Basic Computing 2 tutorial materials, you will find data on scientific articles published between 2004 and 2013 in three top disciplinary journals, *Nature Neuroscience*, *Nature Genetics* and *Ecology Letters*. These data are conatined in three CSV files. We are going to use these data to explore this question.

Load data

Start by reading in the data:

```
data.file <- file.path("citations", "nature_neuroscience.csv")
papers    <- read.csv(data.file, stringsAsFactors = FALSE)
```

Next, take a peek at the data. How large is it?

```
nrow(papers)
ncol(papers)
```

Let's see the first few rows:

```
head(papers)
```

The goal is to test whether papers with longer titles accrue fewer (or perhaps more?) citations than those with shorter titles. The first step is to add another column to the data containing the length of the title for each paper:

```
papers$TitleLength <- nchar(papers$Title)
```

Basic statistics

In the original paper, Letchford *et al* used rank-correlation: rank all the papers according to their title length and the number of citations. If Kendall's τ ("rank correlation") is positive, then longer titles are associated with *more* citations; if τ is negative, longer titles are associated with *fewer* citations. In `R`, you can compute rank correlation using `cor`:

```
k <- cor(papers$TitleLength, papers$Cited.by, method = "kendall")
```

To perform a significance test for correlation, use `cor.test`:

```
k.test <- cor.test(papers$TitleLength, papers$Cited.by, method = "kendall")
```

Does the output of `cor.test` show that the correlation between the ranks is positive or negative? Is this positive or negative correlation significant? You should find that the correlation is opposite of the one reported by Letchford *et al*—longer titles are associated with *more* citations!

Now we are going to examine the data in a different way to test whether this result is robust.

Basic plotting

To plot title length vs. number of citations, we need to learn about plotting in R. To produce a simple scatterplot using the base plotting functions, simply run:

```
plot(papers$TitleLength, papers$Cited.by)
```

The problem with this very simple plot is that it is hard to detect any trend—a few papers have many more citations than the rest, obscuring the data at the bottom of the plot. This suggests that plotting the data on the *logarithmic scale* is a better approach:

```
plot(papers$TitleLength, log10(papers$Cited.by))
```

Question: This is a better plot, but there is one problem with it. What is the problem, and what fix would you suggest? Write your code to fix the problem:

Again, it is hard to see any trend in here. Maybe we should plot the best fitting line and overlay it on top of the graph. To do so, we first need to learn about linear regressions in R.

Linear regression

R was born for statistics — the fact that it is very easy to fit a linear regression in is not surprising! To build a linear model comparing columns x and y in data frame, dat, use `lm`, the "Swiss army knife" for linear regression in R:

```
model <- lm(y ~ x, dat)   # y = a + b*x + error
```

Let's perform a linear regression of the number of citations (on the log-scale) vs. title length. To do so, first create a new column in the data frame containing the counts on the log-scale:

```
papers$LogCits <- log10(papers$Cited.by + 1)
```

Now you can perform a linear regression:

```
model_citations <- lm(LogCits ~ TitleLength, papers)
model_citations          # Gives best-fit line.
summary(model_citations)  # Gives more info.
```

And we can easily add this best-fit line to our plot:

```
plot(papers$TitleLength, papers$LogCits)
abline(model_citations, col = "red", lty = "dotted")
```

Once we add the best-fit line, the positive trend is more clear.

One thing to consider is that this data set spans a decade. Naturally, older papers have had more time to accrue citations. In our models, we should control for this effect. But first, we should explore whether this is an important factor to consider.

First, let's plot the distribution of the number of citations for all the papers:

```
hist(papers$LogCits)
```

You can control the number of histogram bins with the n argument:

```
hist(papers$LogCits, n = 50)
```

Alternatively, estimate the density using density, then plot it:

```
plot(density(papers$LogCits))
```

Next, compare the distributions for papers published in 2004, 2009 and 2013:

```
plot(density(papers$LogCits[papers$Year == 2004]), col = "black")
lines(density(papers$LogCits[papers$Year == 2009]), col = "blue")
lines(density(papers$LogCits[papers$Year == 2013]), col = "red")
```

More recent papers should have fewer citations. Does your plot support this hypothesis? You can account for this in your regression model by incorporating the year of publication into the linear regression:

```
model_citations_better <- lm(LogCits ~ TitleLength + Year, papers)
summary(model_citations_better)
```

Does the regression coefficient (slope) for year confirm that older papers have more citations?

Our new analysis is better than before, but might be even better to have a separate "baseline" for each year. This can be done by converting the "Year" column to a factor:

```r
papers$Year              <- factor(papers$Year)
model_citations_better <- lm(LogCits ~ Year + TitleLength, papers)
summary(model_citations_better)
```

This model has a different baseline for each year, and then title length influences this baseline. In this new model, are longer titles still associated with more citations?

Computing *p*-values using randomization

Kendall's τ takes as input two rankings, x and y, both of the same length, n. It calculates the number of "concordant pairs" (if $x_i > x_j$, then $y_i > y_j$) and the number of "discordant pairs". The final value is

$$\tau = \frac{n_{\text{concordant}} - n_{\text{discordant}}}{\frac{n(n-1)}{2}}$$

If x and y are completely independent, we would expect τ to have a distribution centered at zero. The variance of the "null" distribution of τ depends on the data. It is typically approximated as a normal distribution. If you want to have a stronger result that does not rely on a normality assumption, you can use randomizations to calculate a *p*-value. Simply, compute τ for the actual data, as well as for many "fake" datasets obtained by randomizing the data. Your *p*-value is then the proportion of τ values for the randomized sets that exceed the τ value for the actual data.

Here, we will try to implement this randomization to calculate a *p*-value for papers published in 2006, and then we will compare against the *p*-value obtained from running `cor.test`. To do this, we will use a for-loop for the randomization.

First, subset the data:

```r
dat <- papers[papers$Year == 2006, ]
```

Compute τ from these data:

```r
k <- cor(dat$TitleLength, dat$Cited.by, method = "kendall")
```

Now calculate τ in "fake" data sets by randomly scrambling the citation counts. Begin by doing this for one fake data set:

```r
shuffled_citation_counts <- sample(dat$Cited.by)
k.fake <- cor(dat$TitleLength, shuffled_citation_counts, method = "kendall")
```

Is the value of τ closer to zero in this "shuffled" data set?

To get an accurate *p*-value, we should compute τ for a large number of shuffled data sets. Let's try 1,000 of them. This and similar randomization techniques are known as "bootstrapping".

```
nr       <- 1000        # Number of fake data sets.
k.fake <- rep(0, nr)    # Storage all the "fake" taus.
```

Since this computation involves lots of repetition, a for-loop makes a lot of sense here:

```
for (i in 1:nr){
  shuffled_citation_counts <- sample(dat$Cited.by)
  k.fake[i] <- cor(dat$TitleLength, shuffled_citation_counts,
                   method = "kendall")
}
```

After running this loop, you should have 1,000 correlations calculated from 1,000 fake data sets. You have just generated a "null" distribution for τ. What does this null distribution look like? Try plotting it:

```
hist(k.fake, n = 50)
```

What proportion of the fake data sets have a correlation that exceeds the correlation in the actual data? This is the *p*-value.

```
pvalue <- mean(k.fake >= k)
```

How does your new *p*-value compare to the *p*-value computed by cor.test? Is it smaller or larger?

```
cor.test(dat$TitleLength, dat$Cited.by, method = "kendall")
```

Question: Did you get the same result as the instructor, or your neighbours? If not, why? How could you ensure that your result is more similar, or the same?

Whenever possible, use randomizations rather than relying on classical tests. They are more difficult to implement, and more computationally expensive, but they allow you to avoid making assumptions about your data.

Repeating the analysis for each year

Up until this point, we have only analyzed the citation data for 2006. Does the result we obtained for 2006 hold up in other years? Let's explore this question—we will use a for-loop to repeat the analysis for 2004 to 2013. Let's be smart about designing our code and use a *function* to decompose the problem into parts. The code for the final analysis will look like this:

```
years <- 2004:2013
for (i in years){
  dat <- papers[papers$Year == i, ]
  out <- analyze_citations(dat)
  cat("year:", i, "tau:", out$k, "pvalue:", out$pvalue, "\n")
}
```

The missing piece is the code implementing function `analyze_citations`. You can re-use your code above to write this function.

```
analyze_citations <- function (dat) {
  nr    <- 1000
  k     <- cor(dat$TitleLength, dat$Cited.by, method = "kendall")
  k.fake <- rep(0, nr)
  for (i in 1:nr) {
    k.fake[i] <- cor(dat$TitleLength, sample(dat$Cited.by), method = "kendall")
  }
  return(list(k = k, pvalue = mean(k.fake >= k)))
}
```

Activity: Organizing and running your code

Now we would like to be able to automate the analysis, such that we can repeat it for each journal. This is a good place to pause and introduce how to go about writing programs that are well-organized, easy to write, easy to debug, and easier to reuse.

1. Take the problem, and divide it into smaller tasks (these are the functions).
2. Write the code for each task (function) separately, and make sure it does what it is supposed to do.
3. Document the code so that you can later understand what you did, how you did it, and why.
4. Combine the functions into a master program.

For example, let's say we want to write a program that takes as input the name of files containing citation data. The program should first fit a linear regression model,

```
log(citations + 1) ~ as.factor(Year) + TitleLength
```

then output the coefficient associated with `TitleLength`, and its p-value.

We could split the program into the following tasks:

1. A function to load and prepare the data for a linear regression analysis.
2. A function to run the linear regression analysis.
3. A master code that puts it all together.

Let's begin with the master code—the bulk of the code is a for-loop that repeats the regression analysis for each journal:

```
files <- list.files("citations", full.names = TRUE)
for (i in files) {
  cat("Analyzing data from",i,"\n")
  papers <- load_citation_data(i)
  out    <- fit_citation_model(papers)
  cat("coefficient:", out$estimate, "p-value:", out$pvalue, "\n")
}
```

This code doesn't work yet because you haven't written the functions that are called inside the loop. (What error message to you get when you try to run the code?)

The `load_citation_data` function reads in the data from the CSV file, then prepares the data for the linear regression analysis:

```
load_citation_data <- function (filename) {
  dat <- read.csv(filename, stringsAsFactors = FALSE)
  dat$TitleLength <- nchar(dat$Title)
  dat$LogCits    <- log10(dat$Cited.by + 1)
  dat$Year       <- as.factor(dat$Year)
  return(dat)
}
```

Before continuing, check that it works by running it on one of the CSV files:

```
papers <- load_citation_data("citations/nature_neuroscience.csv")
```

The `fit_citation_model` function fits a linear regression model to the input data, then extracts the quantities from the regression analysis we are most interested in (the "best-fit" slope and the p-value corresponding to "TitleLength").

```
fit_citation_model <- function (papers) {
  model <- lm(LogCits ~ Year + TitleLength, papers)
  terms <- summary(model)$coefficients
  return(list(estimate = terms["TitleLength", "Estimate"],
              pvalue   = terms["TitleLength", "Pr(>|t|)"]))
}
```

Check that this function runs, and does what it is supposed to do:

```
out <- fit_citation_model(papers)
```

Now that you have defined the necessary functions, try running the master code above.

Question: Suppose you download a fourth CSV file containing data on papers from the *American Journal of Human Genetics*. Would any changes need to be made to your R code above to run it on the four citation data sets?

Creating computational notebooks using R Markdown (optional)

Let us change our traditional attitude to the construction of programs: instead of imagining that our main task is to instruct a computer what to do, let us concentrate rather on explaining to humans what we want the computer to do.

Donald E. Knuth, *Literate Programming*, 1984

When doing experiments, it is important to develop the habit of writing down everything you do in a laboratory notebook. That way, when writing your manuscript, responding to queries or discussing progress with your advisor, you can go back to your notes to find exactly what you did, how you did it, and possibly *why* you did it. The same should be true for computational work.

RStudio makes it very easy to build a computational laboratory notebook. First, create a new R Markdown file. (Choose **File > New File > R Markdown** from the RStudio menu bar.)

An R Markdown file is simply a text file. But it is interpreted in a special way that allows the text to be transformed it into a webpage (`.html`) or PDF file. You can use special syntax to render the text in different ways. Here are a few examples of R Markdown syntax:

```
*Italic text* **Bold text**

# Very large header

## Large header

### Smaller header

Unordered and ordered lists:

+ First
+ Second
    + Second 1
    + Second 2

1. This is a
2. Numbered list
```

When rendered as a PDF, the above R Markdown looks like this:

Very large header

Large header

Smaller header

Italic text **Bold text**

Unordered and ordered lists:

- First
- Second
    - Second 1
    - Second 2

1. This is a
2. Numbered list

You can also insert inline code by enclosing it in backticks.

The most important feature of R Markdown is that you can include blocks of code, and they will be interpreted and executed by R. You can therefore combine effectively the code itself with the description of what you are doing.

For example, including a code chunk in your R Markdown file,

```{r hello-world}
cat("Hello world!")
```

will render a document containing both the results and the code that run to generate those results:

```
cat("Hello world!")
# Hello world!
```

If you don't want to run the R code, but just display it, use {r hello-world, eval = FALSE}; if you want to show the output but not the code, use {r hello-world, echo = FALSE}.

You can include plots, tables, and even mathematical equations using LaTeX. In summary, when exploring your data, or describing the methods for your paper, give R Markdown a try!

You can find inspiration in this Boot Camp; the materials for Basic and Advanced Computing were written in R Markdown.